

# A Feature Gradient Approach for Interactive Segmentation

Bharath Sankaran, Mrinal Kalakrishnan, Jeannette Bohg, Stefan Schaal

**Abstract**—The paper is about a really cool idea on how to incorporate the action models to achieve perceptual tasks

## I. INTRODUCTION

## II. RELATED WORK

## III. PROBLEM FORMULATION

Consider a set of rigid objects piled-up on any surface denoted  $T_1, \dots, T_n \in \mathcal{T}$ , where  $n$  is the number of objects. This set of rigid objects can be represented by a graph  $\mathcal{G}$  where each vertex  $g_i \in \mathcal{G}$  corresponds to a node with unique appearance. Here appearance can be defined by a number different attributes combined into a feature vector.

Every object in our set  $\mathcal{T}$  can be represented by a rigid clique in this graph  $\mathcal{G}$  comprising of one or more interconnected nodes  $g_i$ . Edges between rigid cliques are dynamic edges which are removed when two rigid objects are separated. Hence each rigid object corresponds to a set of nodes  $\{g_{i1}, \dots, g_{in}\}$  in a clique  $\mathcal{C}_i$ . A collection of such cliques with

with edges between them represents our scene graph denoted by  $\mathcal{G}$ . We cast the clutter segmentation problem as a graph separation problem where our objective is to pick a set of minimal actions to remove the dynamic edges from the graph and separate the graph into a set of minimal cliques each of which represent a rigid body.

In the initialization phase the dynamic edges between cliques represent the rigid objects that are either touching each other or entangled with each other in a pile of objects.

Given an initial scene graph  $\mathcal{G}$  our objective is to pick the most optimal sequence of actions from a discrete set of actions  $\{a_1, \dots, a_m\} \in \mathcal{A}$  which optimizes the objective of graph separation and reduces the number of actions required to be executed. As the problem of trying to identify the set of actions that optimizes

the graph separation objective for the set of all possible graphs is intractable (given all possible set of graphs that one might encounter in natural scenes). To circumvent this problem we try to learn a mapping from actions  $\mathcal{A}$  to features  $\mathcal{F}$ . These features  $\mathcal{F}$  represent the current state of the environment. Instead of mapping actions to graphs we constraint the problem by trying

to learn a mapping between actions  $\mathcal{A}$  and features  $\mathcal{F}$ . The quality of the features can be evaluated by the reward

observed after executing each action. This learning problem can be cast as a supervised learning problem to classify the features based on action labels. Our learning approach discussed in detail Section V. Since we need to label features, we use Learning from

Demonstration (LfD) to execute actions and label them with user specified rewards. We use a Max Entropy learner as the dimensionality of the feature vector we use is much larger compared to the number of samples we collect.

As the dimensionality of the feature space is incredibly high, we may still not be able to capture the entire variance in the feature space. To account for this variation, we let the learner adapt online to the features. Hence we utilize an online policy gradient styled approach where we optimized parametrized actions with respect to an expected reward by gradient descent. The details of the approach are discussed in Section V.

## IV. LEARNING FROM DEMONSTRATION

In our approach to clutter segmentation, since the learner/robot is not capable of exploring the entire space of all possible adjacency matrices for graph separation, we constraint the problem by exploring in the feature space rather than the space of all possible graphs. To accomplish this we define a set of features over the scene graph and execute actions on them

and observe rewards. Each action is demonstrated by an expert who also labels the actions as good and bad based on the observed reward. So once an expert demonstrates an action, the action is executed on a randomly selected scene graph and the reward is observed. After observing a sequence of such rewards and the corresponding features computed for that specific scene graph,

we feed the features and their corresponding labels (1/0) into a Max Entropy learner to learn the weights for each action. These weights are then used in the online phase to select appropriate actions for a set of observed features. In the online phase the weights are also updated using a policy gradient approach discussed in Section V.

### A. Feature Selection

To learn the utility of actions for a given scene graph  $\mathcal{G}$ , we define a set of features  $\mathcal{F}$  over the scene graph. The features we use are color histograms, textons, grasp templates ([1] TODO: Cite Alex here), push-shape templates and entropy maps. The computation of these features are discussed in detail in the following subsections

This work was supported by the

B. Sankaran, M. Kalakrishnan and S.Schaal are with the Computational Learning and Motor Control Lab, University of Southern California, Los Angeles, CA 90089, bsankara@usc.edu, kalakris@usc.edu  
J.Bohg and S.Schaal are with the Autonomous Motion Department, Max Planck Institute for Intelligent Systems, Tuebingen Germany, jboh@mpi-is.de, sschaal@mpi-is.de

1) *Color Histogram*: To compute a color histogram over a given image patch we convert the patch from rgb to hsv color space and quantize the hue space into 3 bins and saturation space into 6 bins. The number of bins was selected heuristically to satisfy computational efficiency and sparsity constraints. So for each hue bin there are 6 corresponding saturation bins. Hence we quantize the hue-saturation values into 18 bins.

2) *Texon Features*: The idea of textons was first introduced by Olshausen and Field as fundamental micro-structures in generic natural images ([1] TODO: Fix Citation). We utilize the version first introduced by Leung and Malik in ([1] TODO: Fix Citation). Here we convolve the image patch with a series of gabor filters at six evenly spaced orientations at two different pyramid scales.

Once the patch is the image patch is convolved with this series of gabor filters, we take each individual filter response of each image patch and cluster them in an unsupervised manner using k-means clustering. (In our implementation we use k=24 for computational efficiency reasons). Once the responses from the filters are clustered into visual words, we quantize the gabor filter responses corresponding to the computed

visual words using a nearest neighbour approach. This histogram of quantized responses now serves as a texon feature.

3) *Entropy Map*: To quantify the entropy of the local image patch we quantize the colors of the image into RGB histograms. These histograms are normalized and their frequency is computed by summing the bins in each individual channel. The entropy of the image is then computed by summing the shannon entropy of each bin in each channel.

This is computed as shown below:

$$Entropy = \sum_j p_j * \log(p_j)$$

where  $p_j$  is given by (Bin Value)/(Frequency of Channel). Hence this gives us three values for entropy, one for each color channel.

4) *Push Template*: Here we introduce the notion of a push template. A push template is a region in the point cloud which is amenable to pushing by a manipulator. To define such a template we extract holes in images, which are recovered by projecting euclidean clusters on some nominal plane and looking for voids on this plane. Once these voids are detected,

we extract the largest pointcloud surrounding such a void in an image by looking at the total number of points inside an oriented bounding box. The orientations of the bounding box are determined by uniformly discretizing the space of all yaw. The size of the bounding box is constrained by the radius of the size of the gripper of the manipulator in use.

Of all the tested orientations the bounding box with the maximum number of point is selected. We then compute quantized local normals (Fast Point Feature Histograms [1] TODO: Fix Citation) for the pointcloud inside this bounding

box and take the mean of the feature computed over all points.

5) *Grasp Templates*: As a final feature we compute grasp templates which are local heightmaps computed for a given sensor viewpoint and gripper pose. These features were first introduced by Herzog et al in ([1] TODO: Fix Citation). The reader is encouraged to read the paper mentioned to learn about the computation of this specific feature.

Once these features are computed we concatenate these features in to a 112 dimensional feature vector to describe the current scene graph  $mathcal{G}$ . Once we record this feature  $\mathcal{F}$  for a given scene graph, then we execute one of the actions in our action set  $\mathcal{A}$  and label the feature as 1-0 based on the observed reward.

## B. Max Entropy Learning

After number of executions of each action we take the labeled features  $\mathcal{F}_l$  and run them through a Max-Entropy Learner to learn weights on these features. We learn a set of weights  $w_i$  corresponding to each action  $a_1, \dots, a_i \in \mathcal{A}$ . We use L-1 logistic regression as our max entropy learner.

The objective function we try to optimize in our learner is given by:

$$J = \sum_{i=1}^m -\log\left(\frac{1}{1 + \exp^{w^T x_i y_i}}\right) + \lambda ||w||_1$$

This objective function is optimized by the limited memory version of the BFGS algorithm namely, L-BFGS. At test time the labels of a given feature can be determined by

$$y_{test} = \text{sgn}(w^T x_{test})$$

We use L-1 regularized logistic regression as opposed to L-2 regularized logistic regression as it performs better when the dimensionality of the features exceeds that of the number of available training samples. (TODO: Cite that paper that Mrinal cites in his work). We use these learned weights in conjunction with the features observed online by sampling from a gibbs distribution as suggested by ([1] TODO: Fix Citation, Cite Peters and Bagnell)

$$\pi(a|s) \sim \left(\frac{\exp^{\phi(s,a)^T \theta}}{\sum_b \exp^{\phi(s,b)^T \theta}}\right)$$

The action selection and the online adaptation of the weights is discussed in more detail in Section V.

## V. ONLINE ADAPTATION

The online phase of the algorithm we create a scene graph  $\mathcal{G} = \{V, E\}$ , this graph at a selected vertex and update the action parametrization by updating the feature weights by gradient descent in feature space. The online phase of the algorithm has three stages; a supervised classification phase, a data association phase and an online gradient update phase. The three phases are discussed in detail in the sections below.

### A. Supervised Classification

When a scene is presented to the agent (robot) in the form of an rgb image and corresponding point cloud is used to construct a scene graph  $\mathcal{G}$ . We first pre-process the point cloud to look for euclidean clusters supported by some dominant plane in the image. Once these clusters have been detected we project them back on to the image to get the corresponding image patches.

To construct the scene graph we oversegment these image patches in to superpixels using the graph based segmentation approach described in ([1] TODO: Fix Citation). We then construct the scene graph by analyzing the nearest neighbours of each superpixel segment. Every vertex  $V_i$  of the graph  $\mathcal{G}$  would be correspond to a superpixel and the edges  $E_j$  would be be drawn between

superpixels touching each other. From such a representation the notion of rigid and dynamic edges are readily apparent as rigid edges are those edges that connect superpixels on the same object and dynamic edges are those edges that connect superpixels between two objects that are touching or in close proximity to each other. This is illustrated in the figure below Fig 1.



Fig. 1. Illustrates the scene graph construction from a set of objects on a table

Once we generate the scene graph  $\mathcal{G}$ , we compute the set of features described in Section IV for the vertex of maximum degree  $V_{max}$  in the graph. The intuition for only attempting to manipulate the Vertex of maximum degree is that it is connected the most number of other vertices hence manipulating this vertex can impact the graph adjacency matrix more severely than any other node.

Once we compute the features for the max degree vertex  $V_{max}$  we use the weights computed using the max entropy learner in the offline learning phase discussed in Section IV. Using these weights we select the best action to execute by sampling the best action from a gibbs distribution as shown below.

$$a_t = \max(\frac{\exp^{w^T f(x)}}{\sum_b \exp^{w^T f(x)}}$$

Once we generate the action to be executed we execute the action  $a_t$  and observe the corresponding reward  $R_t$ . Here  $t$  is the current time step and  $f(x)$  is the feature function.

### B. Data Association

Once the action is executed we track the vertices of the scene graph  $\mathcal{G}_t$  using optical flow. The positions of the vertices of the graph are updated using the Gunnar Farnebaack Optical Flow algorithm ([1] TODO: Cite Gunnars' ICPR 2000 paper) to obtain the predicted scene graph update  $\tilde{\mathcal{G}}_{t+1}$ . Once the vertices in the graph are updated to get the predicted graph  $\tilde{\mathcal{G}}_{t+1}$ , the predicted graph is matched with the observed scene graph  $\hat{\mathcal{G}}_{t+1}$  to get the actual

scene graph  $\mathcal{G}_{t+1}$ . The graph matching is accomplished by measuring the Bhattacharya distance given by

$$D_B(p, q) = -\log \sum_{x \in \mathcal{X}} \sqrt{p(x)q(x)}$$

between the appearance histograms of the vertices of each graph. Using this distance metric we are able to establish an accurate correspondence between  $\tilde{\mathcal{G}}_{t+1}$  and  $\hat{\mathcal{G}}_{t+1}$ , to get  $\mathcal{G}_{t+1}$ . This correspondence matching is show below in Figure 2.




Fig. 2. Show's the result of the Graph Matching algorithm after the Optical Flow update

Once we get the updated scene graph  $\mathcal{G}_{t+1}$  we perform an online policy gradient update utilizing the expected reward and the observed reward from the updated scene graph.

### C. Online Feature Gradient Update

A policy gradient method is a **reinforcement learning** method that directly parametrized control policy by gradient descent. Since we use discrete actions in our approach; as opposed to performing gradient descent in policy space we perform gradient descent in feature space, thereby adapting the action selection to the current scene being observed. This is accomplished in a method similar to that employed by policy gradient approaches.

Our feature gradient also follows the gradient of the expected return where the feature weights are updated as follows.

$$w_{k+1} = w_k + \alpha_k \nabla_w J(\pi_w)|_{w=w_k}$$

where  $J(\pi_w)$  the gradient on the reward is computed using regular regression

$$\nabla_w = (\Delta W^T \Delta W)^{-1} \Delta W^T \Delta J$$

The mean-subtracted reward return  $\Delta J$  is obtained by subtracting the observed reward from the expected reward. The expected reward is given a function of the features weights computed using the Max Entropy learner in the offline phase. Hence the expected reward is  $J(\hat{w}_k) = W^T f(x)$ . The observed reward is expected reward weighted by the variation in spectral norm of the adjacency matrix. This is given by

$$J(\bar{w}_k) = \frac{1}{(\|A_{t+1}\| - \|A_t\|)} * \beta * J(\hat{w}_k)$$

where  $\beta$  is an appropriate scaling factor. The reason we use the spectral norm is because the spectral norm measures the magnitude of the largest singular value of a matrix. And to measure the similarity between two matrices we can compare their spectral norms. Hence using this logic the gradient update is performed as long as the observed adjacency matrix and the actual adjacency matrix are different.

Once this update is performed a new action is sampled from the gibbs distribution parametrized by the updated feature weights. The actions are sampled until the variation in the spectral norm goes to zero or the feature weights stop updating.

## VI. PERFORMANCE EVALUATION

## VII. CONCLUSION

## REFERENCES

- [1] M. Naghshvar and T. Javidi, "Active sequential hypothesis testing," *IEEE Transactions on Information Theory (to appear)*, March 2012.